

Nicklaus Choo

SOFTWARE ENGINEER • GOOGLE

□ (+1) 669-274-8810 | □ nicklauscyc@gmail.com | □ nicklauscyc | □ nicklaus-choo

Work Experience

Google SOFTWARE ENGINEER | SERVER LOAD BALANCING *SunnyVale, CA, USA*

Aug 2022 - Present

- Designed and Implemented load balancing for ML/server workloads using GPU/TPU/KV-cache utilization. This achieved:
 - 96% faster Time-to-First-Token P90 latency.
 - 60% faster Normalized-Time-Per-Output-Token P90 latency.
 - 32% faster overall throughput for prefix-heavy workloads.
- Led the design and implementation of a distributed load balancing client-server system with partner teams in Canada and Poland. This achieved:
 - Multi-threaded, concurrent, sharded, distributed load balancing client-server resilient to regional outages.
 - Cloud customers can balance service traffic along any utilization dimension.
 - Vastly simpler configuration (50% reduction in configuration programming effort.) for custom load balancing behavior.
- Automated end-to-end testing framework for 70% reduction in overall manual effort.
- 10x speedup for network design process with 15x reduction in memory resources for network topology graph data structures.
- Wrote graph traversal algorithms to efficiently traverse Google's network topology and detect single points of failure.
- Improved automation to re-map 25% of all edge network customers to greatly improve customer cost center allocation.

NetApp SOFTWARE ENGINEER INTERN *SunnyVale, CA, USA*

May 2021 - Aug 2021

- 3x speedup for OS compile-update-reboot time for ONTAP virtual machines.
- Automated ONTAP cloud cluster setup configuration for dynamic load testing.
- 2.7x speedup for WAFL scheduler client I/O latency with minimal slowdown (0.8x) to WAFL scheduler replication operations latency.
- Further optimized 3.5x speedup for client I/O and 1.2x speedup for replication operations with online random forest server load prediction.

ZODAJ FULL STACK SOFTWARE ENGINEER INTERN *Pittsburgh, PA, USA*

May 2020 - Aug 2020

- Devised SMS and Android COVID-19 contact tracing for Senegal in partnership with Senegalese health authorities.
- Designed NoSQL, PostgreSQL databases along with AWS Lambda RESTful API for TCN protocol contact tracing.
- Directed and implemented all permissions and roles within ZODAJ for Amazon AWS databases.
- Created exposure overlap algorithm for n people in a store that runs in $\max(O(k) \text{ or } O(n \log n))$, where k is overlap count.
- Wrote graph algorithms, discrete-time Markov chain SIR epidemic modeling for tracking infections.

Noteworthy Projects

32-bit Kernel

- Created x86 kernel from scratch which supports essential syscalls such as fork/exec/wait, pre-emptive multitasking.
- Wrote device drivers for keyboard, timer, and hardware cursor.
- Code available at github.com/nicklauscyc/small-kernel

Visual Question Answer

- Created a computer vision model to answer questions about images using a Convolutional Neural Network in pytorch
- Code available at github.com/nicklauscyc/vqa

Education

Carnegie Mellon University

Pittsburgh, PA, USA

B.S. IN COMPUTER SCIENCE, CONCENTRATION IN ALGORITHMS & COMPLEXITY WITH UNIVERSITY HONORS

Aug 2018 - May 2022

- School of Computer Science Dean's List, High Honors F19, F20, F21.
- Selected courses:

10-701 Machine Learning (PhD)

15-410 Operating Systems

15-411 Compiler Design (C++)

15-440 Distributed Systems (Java)

15-451 Algorithm Design & Analysis

15-459 Quantum Computation

15-356 Cryptography

15-210 Parallel Data Structures & Algos

15-213 Introduction to Computer Systems

15-251 Great Ideas in Theoretical CS

15-259 Probability & Computing

15-260 Statistics & Computing

21-241 Linear Algebra

21-259 Calculus 3D

21-301 Combinatorics

21-373 Abstract Algebra

80-413 Category Theory

Technical Proficiencies

Languages

- C | C++ | Python | Go | Java | SML
- YAML | Bash | HTML | CSS | JavaScript

Frameworks

Google Kubernetes Engine production server configuration

Borg production cluster turn-up and monitoring

Docker containerized OS build testing, app testing

Terraform infrastructure configuration

Pod release automation / monitoring / observability